

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

sns.set_style('darkgrid')
```

```
In [4]: ## read the dataset

df = pd.read_csv('income_evaluation.csv')
```

```
In [5]: df.head()
```

```
Out[5]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	income
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

```
In [6]: ## check null values in dataset

df.isnull().sum()
```

```
Out[6]: age                0
workclass                0
fnlwgt                  0
education               0
education-num           0
marital-status          0
occupation              0
relationship            0
race                   0
sex                    0
capital-gain            0
capital-loss            0
hours-per-week          0
native-country          0
income                  0
dtype: int64
```

```
In [7]: df.shape
```

```
Out[7]: (32561, 15)
```

```
In [8]: df.columns = ['age', 'workclass', 'final_weight', 'education', 'education_num', 'martial_status', 'occupation', 'relationship', 'race', 'sex',
'capital_gain', 'capital_loss', 'hrs_per_week', 'native_country', 'income']
```

```
In [9]: df.columns
```

```
Out[9]: Index(['age', 'workclass', 'final_weight', 'education', 'education_num',
            'marital_status', 'occupation', 'relationship', 'race', 'sex',
            'capital_gain', 'capital_loss', 'hrs_per_week', 'native_country',
            'income'],
            dtype='object')
```

```
In [10]: df.income.unique()
```

```
Out[10]: array([' <=50K', ' >50K'], dtype=object)
```

```
In [11]: ## converting income columns to 0's and 1's {0 for salary <=50K and 1 for salary >50K}
df['income'] = [1 if value == ' >50K' else 0 for value in df['income'].values]
```

```
In [12]: ## last 5 rows in dataset
```

```
df.tail()
```

```
Out[12]:
```

	age	workclass	final_weight	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss	hrs_per_week	native_country	income
32556	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	0	0	38	United-States	0
32557	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	0	0	40	United-States	1
32558	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	0	0	40	United-States	0
32559	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	0	0	20	United-States	0
32560	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	15024	0	40	United-States	1

```
In [13]: df.workclass.unique()
```

```
Out[13]: array([' State-gov', ' Self-emp-not-inc', ' Private', ' Federal-gov',
            ' Local-gov', ' ?', ' Self-emp-inc', ' Without-pay',
            ' Never-worked'], dtype=object)
```

```
In [14]: ## removing ? from column workclass
```

```
df['workclass'] = np.where(df.workclass == '?', np.nan, df['workclass'])
```

```
In [15]: df.dropna(axis=0, inplace=True)
```

```
In [16]: workclass_label = {v:k for k, v in enumerate(df.workclass.unique())}
```

```
In [17]: workclass_label
```

```
Out[17]: {' State-gov': 0,
            ' Self-emp-not-inc': 1,
            ' Private': 2,
            ' Federal-gov': 3,
            ' Local-gov': 4,
            ' Self-emp-inc': 5,
            ' Without-pay': 6,
            ' Never-worked': 7}
```

```
In [22]: df.workclass = df.workclass.map(workclass_label)
```

```
In [23]: df.education.unique()
```

```
Out[23]: array([' Bachelors', ' HS-grad', ' 11th', ' Masters', ' 9th',  
              ' Some-college', ' Assoc-acdm', ' Assoc-voc', ' 7th-8th',  
              ' Doctorate', ' Prof-school', ' 5th-6th', ' 10th', ' Preschool',  
              ' 12th', ' 1st-4th'], dtype=object)
```

```
In [24]: education_label = {v:k for k, v in enumerate(df.education.unique())}
```


```
In [25]: education_label
```

```
Out[25]: {' Bachelors': 0,  
          ' HS-grad': 1,  
          ' 11th': 2,  
          ' Masters': 3,  
          ' 9th': 4,  
          ' Some-college': 5,  
          ' Assoc-acdm': 6,  
          ' Assoc-voc': 7,  
          ' 7th-8th': 8,  
          ' Doctorate': 9,  
          ' Prof-school': 10,  
          ' 5th-6th': 11,  
          ' 10th': 12,  
          ' Preschool': 13,  
          ' 12th': 14,  
          ' 1st-4th': 15}
```

```
In [26]: df.education = df.education.map(education_label)
```

```
In [27]: plt.figure(figsize=(10, 4))  
sns.countplot(x='income', data=df, hue='sex');
```

```
In [31]: plt.figure(figsize=(10, 5))  
plt.xticks(rotation=75)  
sns.barplot(x=native_country.index, y=np.log(native_country))  
plt.ylabel('Log of Count')  
plt.show()
```



```
In [32]: df.native_country = np.where(df.native_country == ' ?', np.nan, df['native_country'])
```

```
In [33]: df.dropna(axis=0, inplace=True)
```

```
In [34]: native_country = {v:k for k, v in enumerate(df.native_country.unique())}
```

```
In [35]: df.native_country = df.native_country.map(native_country)
```

```
In [36]: df.head()
```

```
Out[36]:
```

	age	workclass	final_weight	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss	hrs_per_week	native_country	income
0	39	0	77516	0	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	0	0
1	50	1	83311	0	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	0	0
2	38	2	215646	1	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	0	0
3	53	2	234721	2	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	0	0
4	28	2	338409	0	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	1	0

```
In [37]: df.marital_status.unique()
```

```
Out[37]: array([' Never-married', ' Married-civ-spouse', ' Divorced',  
             ' Married-spouse-absent', ' Separated', ' Married-AF-spouse',  
             ' Widowed'], dtype=object)
```

```
In [38]: marital_label = {v:k for k, v in enumerate(df.martial_status.unique())}
```

```
In [39]: df.martial_status = df.martial_status.map(marital_label)
```

```
In [40]: df.occupation.unique()
```

```
Out[40]: array([' Adm-clerical', ' Exec-managerial', ' Handlers-cleaners',  
             ' Prof-specialty', ' Other-service', ' Sales', ' Transport-moving',  
             ' Farming-fishing', ' Machine-op-inspct', ' Tech-support',  
             ' Craft-repair', ' Protective-serv', ' Armed-Forces',  
             ' Priv-house-serv', ' ?'], dtype=object)
```

```
In [41]: df.occupation = np.where(df.occupation == ' ?', np.nan, df['occupation'])
```

```
In [42]: df.dropna(axis=0, inplace=True)
```

```
In [43]: occ_label = {v:k for k, v in enumerate(df.occupation.unique())}
```

```
In [44]: df.occupation = df.occupation.map(occ_label)
```

```
In [45]: df.relationship.unique()
```

```
Out[45]: array([' Not-in-family', ' Husband', ' Wife', ' Own-child', ' Unmarried',  
             ' Other-relative'], dtype=object)
```

```
In [46]: relationship_label = {v:k for k, v in enumerate(df.relationship.unique())}
```

```
In [47]: df.relationship = df.relationship.map(relationship_label)
```

```
In [48]: df.head()
```

```
Out[48]:
```

	age	workclass	final_weight	education	education_num	martial_status	occupation	relationship	race	sex	capital_gain	capital_loss	hrs_per_week	native_country	income
0	39	0	77516	0	13	0	0	0	White	Male	2174	0	40	0	0
1	50	1	83311	0	13	1	1	1	White	Male	0	0	13	0	0
2	38	2	215646	1	9	2	2	0	White	Male	0	0	40	0	0
3	53	2	234721	2	7	1	2	1	Black	Male	0	0	40	0	0
4	28	2	338409	0	13	1	3	2	Black	Female	0	0	40	1	0

```
In [49]: df.sex = np.where(df.sex == ' Male', 1, 0)
```

```
In [50]: df.race.unique()
```

```
Out[50]: array([' White', ' Black', ' Asian-Pac-Islander', ' Amer-Indian-Eskimo',  
             ' Other'], dtype=object)
```

```
In [51]: race_label = {v:k for k, v in enumerate(df.race.unique())}
```

```
In [52]: race_label

Out[52]: {' White': 0,
          ' Black': 1,
          ' Asian-Pac-Islander': 2,
          ' Amer-Indian-Eskimo': 3,
          ' Other': 4}

In [53]: df.race = df.race.map(race_label)


In [54]: df.head()


Out[54]:
```

	age	workclass	final_weight	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss	hrs_per_week	native_country	income
0	39	0	77516	0	13	0	0	0	0	1	2174	0	40	0	0
1	50	1	83311	0	13	1	1	1	0	1	0	0	13	0	0
2	38	2	215646	1	9	2	2	0	0	1	0	0	40	0	0
3	53	2	234721	2	7	1	2	1	1	1	0	0	40	0	0
4	28	2	338409	0	13	1	3	2	1	0	0	0	40	1	0

```
In [56]: plt.figure(figsize=(15, 15))
sns.heatmap(df.corr(), annot=True)
plt.show()
```










v  
fin  
e  
educat  
mariti  
oc  
rel  
cap  
cap  
hrs\_f  
native

```
In [57]: df.head()
```


Out[57]:

	age	workclass	final_weight	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss	hrs_per_week	native_country	income
0	39	0	77516	0	13	0	0	0	0	1	2174	0	40	0	0
1	50	1	83311	0	13	1	1	1	0	1	0	0	13	0	0
2	38	2	215646	1	9	2	2	0	0	1	0	0	40	0	0
3	53	2	234721	2	7	1	2	1	1	1	0	0	40	0	0
4	28	2	338409	0	13	1	3	2	1	0	0	0	40	1	0

```
In [59]: for feature in ['final_weight', 'capital_gain']:  
    plt.figure(figsize=(10, 5))  
    sns.boxplot(df[feature])  
    plt.show()
```









In [60]: `df.describe()`

Out[60]:

	age	workclass	final_weight	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss	hrs_per_wee
count	30162.000000	30162.000000	3.016200e+04	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000	30162.000000
mean	38.437902	2.109343	1.897938e+05	3.368842	10.121312	1.075061	4.615609	1.523971	0.211823	0.675685	1092.007858	88.372489	40.93123
std	13.134665	0.934785	1.056530e+05	3.404320	2.549995	1.217557	3.432195	1.431980	0.612461	0.468126	7406.346497	404.298370	11.97998
min	17.000000	0.000000	1.376900e+04	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000
25%	28.000000	2.000000	1.176272e+05	1.000000	9.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	40.000000
50%	37.000000	2.000000	1.784250e+05	2.000000	10.000000	1.000000	4.000000	1.000000	0.000000	1.000000	0.000000	0.000000	40.000000
75%	47.000000	2.000000	2.376285e+05	5.000000	13.000000	1.000000	8.000000	3.000000	0.000000	1.000000	0.000000	0.000000	45.000000
max	90.000000	6.000000	1.484705e+06	15.000000	16.000000	6.000000	13.000000	5.000000	4.000000	1.000000	99999.000000	4356.000000	99.000000

In [61]: `from sklearn.model_selection import train_test_split`  
`from sklearn.preprocessing import StandardScaler`

```
In [62]: X = df.iloc[:, :-1]
y = df.iloc[:, -1]
```

```
In [63]: X = X.values
y = y.values

sc = StandardScaler()
sc.fit_transform(X)
```

```
Out[63]: array([[ 0.04279571, -2.25653747, -1.0627216 , ..., -0.21858598,
                -0.07773411, -0.23923831],
                [ 0.88028814, -1.18675527, -1.00787131, ..., -0.21858598,
                -2.3315307 , -0.23923831],
                [-0.03333996, -0.11697307,  0.24469349, ..., -0.21858598,
                -0.07773411, -0.23923831],
                ...,
                [ 1.48937355, -0.11697307, -0.3585745 , ..., -0.21858598,
                -0.07773411, -0.23923831],
                [-1.25151078, -0.11697307,  0.11070545, ..., -0.21858598,
                -1.74721307, -0.23923831],
                [ 1.0325595 ,  3.09237353,  0.92884082, ..., -0.21858598,
                -0.07773411, -0.23923831]])
```

```
In [64]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [65]: from sklearn.linear_model import LogisticRegression
```

```
In [66]: lg = LogisticRegression()
```

```
In [67]: lg.fit(X_train, y_train)
```

C:\Users\amiglani\AppData\Local\anaconda3\Lib\site-packages\sklearn\linear\_model\\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

```
Out[67]: LogisticRegression ⓘ ?
LogisticRegression()
```

```
In [68]: pred = lg.predict(X_test)
```

```
In [69]: from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
In [70]: print("-----Classification Report-----")
print(classification_report(y_test, pred))

print("-----Accuracy Score-----")
print(accuracy_score(y_test, pred))

print("-----Confusion Matrix-----")
plt.figure(figsize=(10,5))
```

```
sns.heatmap(confusion_matrix(y_test, pred), annot=True);
```

```
-----Classification Report-----
              precision    recall  f1-score   support

     0       0.79       0.97       0.87       6767
     1       0.70       0.23       0.35       2282

 accuracy          0.78       9049
 macro avg       0.74       0.60       0.61       9049
 weighted avg    0.77       0.78       0.74       9049

-----Accuracy Score-----
0.780970272958338

-----Confusion Matrix-----
```

```
In [72]: from sklearn.metrics import accuracy_score
```

```
# Fit a basic model and check accuracy
from sklearn.tree import DecisionTreeClassifier

model = DecisionTreeClassifier(random_state=42)
model.fit(X_train, y_train)

# Evaluate on training and test sets
train_accuracy = accuracy_score(y_train, model.predict(X_train))
test_accuracy = accuracy_score(y_test, model.predict(X_test))

print(f"Training Accuracy: {train_accuracy}")
print(f"Test Accuracy: {test_accuracy}")
```

```
Training Accuracy: 1.0
Test Accuracy: 0.8060559177809703
```

```
In [74]: from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from xgboost import XGBClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

```
models = {
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "SVM": SVC(),
    "XGBoost": XGBClassifier()
}

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"{name} Accuracy: {accuracy}")
```

```
Decision Tree Accuracy: 0.8081555973035695
Random Forest Accuracy: 0.8541275279036358
SVM Accuracy: 0.7810807824068958
XGBoost Accuracy: 0.8697093601502929
```

```
In [82]: from sklearn.model_selection import cross_val_score
```

```
from sklearn.ensemble import RandomForestClassifier

# Initialize the RandomForestClassifier
rf_classifier = RandomForestClassifier()

# Perform cross-validation
cv_scores = cross_val_score(rf_classifier, X_train, y_train, cv=5)

# Output the results
print(f'Cross-validation scores: {cv_scores}')
print(f'Mean CV score: {cv_scores.mean()}')
```

Cross-validation scores: [0.85578972 0.8598153 0.84915937 0.85054477 0.85315017]  
Mean CV score: 0.8536918633640213

In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js